Udacity's Deep Reinforcement Learning Nanodegree Project 3 Report: Unity's Tennis Environment solved by DDPG

December 25, 2018

1 Introduction

This project involves solving the environment of a double-jointed arm following a target simulated in the Unity MLAgents environment. This environment has a continuous action space which makes it difficult to apply DQN approaches. Thus, we will investigate a more appropriate algorithm called Deep Deterministic Policy Gradient (DDPG) [1, 2].

The double-jointed arm moves about in the 3D space trying to stay in the moving target region. The agent is able to observe 33 dynamic variables about the arm such as velocity, position, angular velocity, etc. The actions are 4 continuous variables representing the torques at the joints of the double-jointed arm.

2 Learning Algorithm

The Deep Deterministic Policy Gradient (DDPG) algorithm [1, 2] is implemented in this project:

1: Initialize replay memory D with capacity N
2: Initialize critic network $\hat{q}$ with random weights w^q
3: Initialize actor network μ with weights w^μ
4: Initialize target critic weights $w^{q-} \leftarrow w^q$
5: Initialize target actor weights $w^{\mu-} \leftarrow w^\mu$
6: **for** the episode $e \leftarrow 1$ to M **do**
7:     Initialize a random process $\mathcal{N}$ for action exploration
8:     Receive initial input state S
9:     **for** time step $t \leftarrow 1$ to T **do**
10:         Choose action $A = \mu(S|w^\mu) + \mathcal{N}$
11:         Take action A, observe reward R, and next input frame S'
12:         Store experience (S, A, R, S') in replay memory D
13:         $S \leftarrow S'$

14:        Obtain minibatch of tuples (s_j, a_j, r_j, s_{j+1}) from D of size K.

15:        Set target $y_j = r_j + \gamma \hat{q}(s_{j+1}, \mu(S', w^{\mu-}), w^{q-})$

16:        Update w^q: $\Delta w^q = -\alpha \frac{1}{N} \sum_j (y_j - \hat{q}(s_j, a_j, w^q)) \nabla_{w^q} \hat{q}(s_j, a_j, w^q)$

17:        Update w^μ with policy gradient:

$$\nabla_{w^\mu} J \approx \frac{1}{N} \sum_i \nabla_a \hat{q}(s_i, a\mu(s_i)) \nabla_{w^\mu} \mu(s_i | w^\mu)$$

18:        Soft update w^{q-}: $w^{q-} \leftarrow (1-\tau)w^{q-} + \tau w^q$

19:        Soft update $w^{\mu-}$: $w^{\mu-} \leftarrow (1-\tau)w^{\mu-} + \tau w^\mu$

The hyperparameters were set to the following: buffer size $N = 10^6$, discount factor $\gamma = 0.99$, batch size $K = 256$, soft update parameter $\tau = 10^{-3}$, learning rate $\alpha = 1 \times 10^{-4}$ for the actor and learning rate $\alpha = 3 \times 10^{-4}$ for the critic.

The critic network used to represent μ is a fully connected neural network with an input layer with 33 input nodes. The second layer is fully connected to the first input layer with 256 nodes and is concatenated with a series of new input nodes for the action. The third layer has 128 nodes. The final layer has a single node resenting the action value for the state and action.

The actor network used to represent $\hat{q}$ is a fully connected neural network with an input layer of 33 nodes, a hidden layer of 256 nodes, and the final output layer with 4 nodes (one for each action value). The activation used for the hidden layer was a ReLU function and a tanh function for the final layer.

3   Results and Plots of Rewards

4   Ideas for Future Work

I am curious to see if learning could be improved by varying noise. I would expect that it is beneficial to have lots of noise at the beginning of training and then less noise at later stages — similar to decaying epsilon over time when using epsilon-greedy.

References

[1] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, sep 2015.

[2] Udacity. Deep Reinforcement Learning Nanodegree Course Material, 2018.